

Jane Communication Automation

Handoff & Migration Guide

Prepared for: Gar (receiving engineer) and the Schilmoeller & Schoenfield team

Prepared by: PRIME (RJ Tohay — rj@primelive.ai; early phases by Nate Geraldez — nate@primelive.ai)

Date: 2026-06-04

Goal: Migrate the system from PRIME's Railway environment to the client's own server, and hand over day-to-day operation.

1. What this system is

An AI-assisted email platform for the firm's Microsoft 365 mailbox (`jane@schilcpa.com`). It polls the mailbox, categorizes every incoming email with AI (Anthropic Claude), assigns a triage tier, drafts a suggested reply from the firm's knowledge base, and gives staff a review UI. **Nothing is ever sent without human approval** (two independent safety gates — see §6.2). It also supports composing new outbound email, attachments in both directions, spam/delete synced back to Outlook, per-client saved folders, full-text search, an escalation queue for Jane, and a complete audit trail.

- **Backend:** Python 3.12 / FastAPI / SQLAlchemy / Alembic — `backend/`
- **Frontend:** Next.js 14 / TypeScript / Tailwind — `frontend/`
- **Database:** PostgreSQL 16 (13 tables, migrations 001–017)
- **In-app user guide:** the app ships with an interactive `/tutorials` page — that is the end-user manual for Jane and staff.

The technical deep-dive (architecture, full API list, schema) is in the repository **README.md**. This document covers what the README doesn't: access, migration steps, operations, and the sharp edges.

2. Repositories & git workflow

Remote	URL	Role
origin (fan-out)	<code>github.com/rheynardjan/schillerCPA</code> and <code>github.com/Janedev99/Communications-Automation</code>	One <code>git push</code> updates both
Client's copy	<code>github.com/Janedev99/Communications-Automation</code>	The authoritative copy the client owns. Migrate from here.

- **Branches:** `master` = production, `development` = integration. Feature work happens on `FEAT/` , `FIX/` , `CHORE/` , `DOCS/` branches off `development` ; merges are `--no-ff` .
- `master` is only updated from `development` with the firm's approval.
- After handoff, Gar can simplify to whatever workflow he prefers — the only hard rule worth keeping: **don't push straight to `master` without testing**, because production migrations run automatically on deploy (§4).

3. Access checklist — what Gar needs to receive

#	Item	From	Notes
1	GitHub access to <code>Janedev99/Communications-Automation</code>	Jane / Craig	Already in the client's org
2	Railway project access (or a fresh DB dump at cutover)	PRIME (RJ)	Needed once, for the data migration
3	Azure app registration credentials (<code>MSGRAPH_TENANT_ID</code> , <code>MSGRAPH_CLIENT_ID</code> , <code>MSGRAPH_CLIENT_SECRET</code>)	PRIME (RJ) → rotate after	App permissions: <code>Mail.Read</code> , <code>Mail.Send</code> , <code>Mail.ReadWrite</code> (application, admin-consented). Rotate the client secret after migration so PRIME no longer holds a working credential. The registration lives in the firm's Azure

			tenant — it keeps working from any host.
4	Anthropic API key	Firm should create its own	console.anthropic.com → API Keys. Don't inherit PRIME's key — billing and rate limits should be the firm's.
5	APP_SECRET_KEY	Generate fresh	<code>python -c "import secrets; print(secrets.token_hex(32))"</code> — do NOT reuse PRIME's
6	RunPod account	Gar already owns it	Optional — only if the self-hosted LLM path is ever revived (\$6.5). The pod used during development is on Gar's shared account.
7	Slack webhook URL (escalation notifications)	Firm's Slack admin	Optional — SLACK_WEBHOOK_URL
8	Admin login for the app	Set via ADMIN_EMAIL/ADMIN_PASSWORD env + <code>seed_admin.py</code>	Existing user accounts come across with the DB dump

4. Current production environment (what you're migrating FROM)

Hosted on **Railway**, three services:

Service	Build	Start
Backend	backend/Dockerfile	<code>sh -c 'alembic upgrade head && uvicorn app.main:app --host 0.0.0.0 --port \${PORT:-8000} --workers 1'</code>
Frontend	frontend/Dockerfile (multi-stage, Next standalone)	<code>node server.js</code>
Postgres	Railway plugin	DATABASE_URL injected

Three things to internalize about this setup:

1. **Migrations run on every backend deploy** (`alembic upgrade head` in the start command). Convenient, but a bad migration can block boot — the health check (`/health` , 300s timeout) will catch it.
2. `NEXT_PUBLIC_API_URL` is a **Docker BUILD ARG**, not a runtime env var. It's inlined into the JS bundle at build time. On Railway it's set under *Build* settings; on any new host it must be passed to `docker build` . Getting this wrong produces a frontend that builds fine but talks to the wrong backend.
3. **Deploy drift has happened before**. In May 2026 production silently ran a stale build while master had moved on (migration lag was the symptom). After every deploy, verify the running version actually changed — check a recently-shipped feature or the migration level (`SELECT version_num FROM alembic_version;` should read `017` as of this writing).

5. Migration runbook (Railway → client server)

5.1 Recommended target: docker-compose

The repo ships a production-shaped `docker-compose.yml` : Postgres 16 + backend + frontend + **Caddy** (reverse proxy with automatic TLS on ports 80/443). This is the lowest-friction self-hosting path.

Alternative: run the two Dockerfiles under any orchestrator, or bare-metal with systemd — the only invariants are PostgreSQL 16+, a single backend worker, and the env vars.

5.2 Step-by-step

Prepare (no downtime):

1. Provision the host (Docker + docker-compose). Open 80/443.
2. Clone `Janedev99/Communications-Automation` .
3. `cp .env.example .env` and fill in everything (\$3 items). Production checklist:
 - `APP_ENV=production`
 - `APP_SECRET_KEY` = freshly generated (boot **fails** on the dev default in production)
 - `LLM_PROVIDER=anthropic` , `ANTHROPIC_API_KEY` = the firm's key
 - `EMAIL_PROVIDER=msggraph` + the four `MSGGRAPH_*` values
 - `SHADOW_MODE=true` (keep auto-send off — see §6.2)
 - `CORS_ORIGINS` = the real frontend URL only (no localhost)
 - `NEXT_PUBLIC_API_URL` = the real backend URL (build arg!)
 - `TRUSTED_PROXIES` = Caddy's container IP for accurate client IPs (blank is safe)

4. Build images. Do a dry-run boot against an EMPTY local database first: backend should apply migrations 001–017 and pass `/health` . This proves the stack works before touching real data.

Cutover (brief downtime, do it outside business hours):

5. **Stop the Railway backend service first.** The poller marks emails processed as it ingests them — two pollers against two databases would each process new mail independently. Stop it, then dump.
6. Dump production data: `pg_dump "$RAILWAY_DATABASE_URL" --no-owner --no-privileges -Fc -f jane_prod.dump`
7. Restore into the new Postgres: `pg_restore --no-owner --no-privileges -d "$NEW_DATABASE_URL" jane_prod.dump`
8. Start the new stack. The backend's `alembic upgrade head` will no-op (the dump already contains schema + `alembic_version`).
9. **Verify** (15 minutes, with Jane or Sara available):
 - `GET /health` returns ok; login works with an existing account
 - Dashboard shows the historical threads/stats (proves the data came across)
 - Send a test email TO the mailbox from a personal account → it appears in the inbox within ~60s with a category and a draft (proves poller + AI)
 - Open an old message with an attachment → download works (proves Graph credentials)
 - Reply-send the test thread WITH an attachment (Jane approves) → recipient receives it; sent bubble shows the attachment chip (proves the send path)
 - Check Settings → Integrations: all cards green
10. Update Jane's browser bookmark (and staff bookmarks) to the new URL. The app is accessed by bookmark — there is no other entry point to update.
11. Decommission: delete the Railway services, **rotate the Azure client secret**, and PRIME removes its copies of all credentials.

5.3 Rollback

Until step 11, the Railway environment is intact (just stopped). Rollback = stop new stack, restart Railway backend, revert bookmarks. The only divergence risk is mail that arrived during cutover — the poller backfills on restart since it dedupes by Message-ID, so nothing is lost either way.

6. Operations guide

6.1 Background jobs (all in-process — hence single worker)

Job	Interval	What it does
Email poller	60s (EMAIL_POLL_INTERVAL_SECONDS)	Fetch → dedupe → categorize → tier → draft → escalate
Session cleanup	hourly	Purges expired login sessions
RunPod watchdog	60s	Idle-stops the GPU pod (no-op when RUNPOD_POD_ID empty)

Never run more than one backend worker/replica — each would run its own poller.

6.2 The two auto-send gates (IMPORTANT)

Auto-sending of T1 drafts only happens if **both** are open:

1. `SHADOW_MODE=false` (environment variable)
2. `auto_send_enabled=true` (Settings UI → system_settings table)

Production today: `SHADOW_MODE=true` → nothing auto-sends, every email is human-approved. **Do not change without Jane's explicit sign-off.** Drafts still generate in shadow mode — that's by design.

6.3 AI budget

`DAILY_TOKEN_BUDGET` (default 1,000,000 tokens/day, resets UTC midnight) caps all AI spend. Exhausted → categorization falls back to the keyword rules engine; drafting pauses until reset. The dashboard shows usage. Bulk reprocessing scripts can burn the full budget in one run — plan those.

6.4 Routine admin tasks

- **Signature:** Settings UI → stored in `system_settings.draft_signature` ; appended to every draft/composed email in code.
- **Triage rules:** Settings → Triage Rules — per-category T1 eligibility + confidence thresholds (complaint/urgent/promotional are locked server-side).
- **Knowledge base:** Knowledge page — the content the AI drafts from. Keep it current; it matters more than prompt tweaks.
- **Users:** Settings (admin) — staff vs admin roles.
- **Audit/Export:** Audit Log page (admin); `GET /api/v1/emails/export` streams JSON-lines for compliance (rate-limited).
- **Backups:** Railway handled this implicitly. On self-hosting, schedule `pg_dump` (daily, retained ≥30 days). The database is the entire system state — email bodies, drafts, KB, audit log. Attachment binaries are NOT in the DB (streamed on demand from the mailbox), so DB backups stay small.

6.5 RunPod (dormant, optional)

The system can run drafting on a self-hosted GPU pod (RunPod) instead of Anthropic — built before the 2026-05-14 decision to standardize on Anthropic for email. The orchestration (auto-start, idle-stop, daily cost cap, Claude fallback) is fully functional but **disabled in production** (`RUNPOD_POD_ID` empty). The dev pod lives on Gar's shared RunPod account (~\$2.99/hr H100 when running). Ignore unless the firm revisits self-hosting; everything is documented in `.env.example`.

7. Known issues & sharp edges

#	Issue	Impact / workaround
1	<code>scripts/seed_demo.py</code> writes to whatever <code>DATABASE_URL</code> points at .	Running it with a production <code>.env</code> inserts demo threads into prod. Only run against a dev database.
2	IMAP provider lacks <code>move_message</code> (<code>NotImplementedError</code>).	Spam/Delete mailbox mirroring works on MS Graph only. Only matters if <code>EMAIL_PROVIDER</code> ever flips to <code>imap</code> .
3	Attachment downloads resolve from the live mailbox .	If a message is hard-deleted from Outlook/Exchange, its attachments 404 in the app (metadata stays, binary is gone). By design — no binary storage.
4	<code>NEXT_PUBLIC_API_URL</code> is build-time .	Changing the backend URL requires rebuilding the frontend image, not just restarting it.
5	Deploy drift precedent .	Verify each deploy actually went live (§4).
6	Mobile layout is not yet refined .	Jane accesses via phone browser; it works but isn't optimized. On the roadmap (§8).
7	Older drafts predating the signature feature lack the configured signature.	Deliberate — firm chose not to regenerate them.
8	Single-worker constraint (§6.1).	Scaling out requires extracting the poller first.

8. Open roadmap (not blocking handoff)

- Mobile view refinement** — committed verbally in the 2026-05-27 client meeting; not yet built.
- Direction discussed with the firm: evolve toward Jane's primary mail client (folders parity, richer search). Discussion stage only.
- "What's New" release-notes surface — designed, not built.
- In-app practice mode for the Compose flow on the tutorials page.

9. Quick reference

- Run tests:** `cd backend && venv/Scripts/python -m pytest tests/ -q` — 356 tests, in-memory SQLite, all providers mocked, zero real sends. Safe anywhere, run before every deploy.
- DB schema / API reference:** `README.md` §Database Schema + live OpenAPI at `<backend>/docs`.
- Local dev:** backend on port **8001** (`uvicorn app.main:app --port 8001`), frontend `npm run dev` on 3000.
- Migration level today:** `017` (`seed draft signature`).
- Contacts during transition:** RJ Tohay — rj@primelive.ai (current development); Nate Geraldez — nate@primelive.ai (original prototype, stages 1–2).

Generated 2026-06-04. The `README` and `.env.example` in the repository are kept current and take precedence over this snapshot if they ever disagree.